

Bankruptcy.R

Sharada Sarangan

2023-03-16

```
# Final Project
```

```
# Set working directory to current required folder
```

```
library(class)
library(MASS)
library(e1071)
library(tree)
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```
library(smotefamily)
library(boot)
library(reshape2)
library(ggplot2)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-6
```

```
options(scipen=999)
```

```
# Part 1
```

```
# EXPLORATORY EDA AND SETTING UP TRAIN AND TEST SET
```

```
dat <- read.csv("bankruptcy_data.csv")
```

```
table(dat$bankrupt)
```

```
##
```

```
##      0      1
```

```
## 6599  220
```

```

# we drop the repeated and redundant columns to prevent multi-collinearity
# based on domain knowledge, since several measures correspond to the same thing
# but have been recoded in different ways for accounting purposes.
dat <- dat[,c(1,7:16,19:22, 25:38,40,41,45:77,83, 86, 87:92, 94)]
# str(dat)

```

```

# setting up training and testing set
# separating the dependent variable
yes <- subset(dat, dat$bankrupt == 1)
no <- subset(dat, dat$bankrupt == 0)

```

```

# setting seed
set.seed(112233)

```

```

# Randomly selecting 1550 rows from both groups of
# of numbers (1-->50 and 1-->1500)
test.yes <- sample(1:nrow(yes),50)
test.no <- sample(1:nrow(no),1500)

```

```

# separating our testing set
dat.test <- rbind(yes[test.yes,],no[test.no,])
# checking our train set
table(dat.test$bankrupt)

```

```

##
##      0      1
## 1500    50

```

```

# creating the training set
newyes <- yes[-test.yes,]
newno <- no[-test.no,]

# our base train set
train_og <- rbind(newyes,newno)

# Check results
table(train_og$bankrupt)

```

```

##
##      0      1
## 5099   170

```

```

# bootstrapped train set (with replacement)
train.newyes <- sample(1:nrow(newyes),500, replace=TRUE)
train.newno <- sample(1:nrow(newno),3000)

# our base train set
train_boot <- rbind(newyes[train.newyes,],newno[train.newno,])

# Check results
table(train_boot$bankrupt)

```

```
##
##      0      1
## 3000  500

# ADAS bootstrapping - Generate synthetic positive instances using ADASYN algorithm
adas_train <- ADAS(train_og,train_og$bankrupt,K=5)
train_adas <- adas_train$data

# dropping the last class column since its the same as bankrupt
train_adas <- train_adas[,c(1:73)]

# lets look at our synthetic values
# adas_train$syn_data

# checking our distribution
table(train_adas$bankrupt)
```

```
##
##      0      1
## 5099 5152
```

```
# removing unnecessary variables
rm(yes, no, test.yes, test.no, newyes, newno, train.newyes,
    train.newno)

# we run an initial log_reg to identify significant variables
glm_var <- glm(bankrupt ~ ., data=train_og, family='binomial')
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(glm_var)
```

```
##
## Call:
## glm(formula = bankrupt ~ ., family = "binomial", data = train_og)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
##    -8.49     0.00     0.00     0.00     8.49
##
## Coefficients:
##              Estimate              Std. Error
## (Intercept) 1883979708516677451776.0000000    1699064189984867.7500000
## X6          -465183775008882163712.0000000    15916642620196.4003906
## X7           388119709635854991360.0000000    13294282568795.6269531
## X8           283983102233256640.0000000      5713580607.6899691
## X9          -222377426493851041792.0000000    7609152820787.1191406
## X10         137351940171298880.0000000      2991145588.2508001
```

## X11	-2781.8383180	0.0003187
## X12	6406.0329957	0.0003657
## X13	-11674391317988228.0000000	199034816.6428339
## X14	-1128938.4341195	0.0086069
## X15	-903101588679473.1250000	7154954.4825070
## X18	-1590748969963942.5000000	48770603.0520286
## X19	-6469232962841985.0000000	66792839.3558250
## X20	1176030986828447.0000000	107148022.4485826
## X21	2272229.4327178	0.0581123
## X24	-407791986138473.4375000	69223260.9213606
## X25	-1992878399596448.5000000	108281472.1661287
## X26	12869912493968414.0000000	1253497879.9802492
## X27	-11458159198611292.0000000	1250207134.1004519
## X28	162820692948075.0937500	81826413.3755742
## X29	-2761.4022618	0.0003311
## X30	388483.4184615	0.0082377
## X31	1484647239702127.2500000	99098898.8917462
## X32	2226928032426705.2500000	111528249.2206729
## X33	-1246098.0609360	0.0245209
## X34	51066.5822066	0.0035889
## X35	5273774632004118.0000000	121049302.7006431
## X36	237993.2494538	0.0067400
## X37	3390971075940789.5000000	87957268.2152067
## X39	-1059782924336249.1250000	40519475.6397451
## X40	9559741624139022.0000000	293710759.3670629
## X44	-3042218055542296.0000000	336243559.8469479
## X45	-838239631906503.0000000	21336378.3386176
## X46	-889060.6917348	0.0038217
## X47	-160828.6135705	0.0046055
## X48	-2248.4327626	0.0003031
## X49	7151.8508011	0.0004539
## X50	-46762232216170.1875000	57622024.0116544
## X51	1067902.9366004	0.0205171
## X52	646895347379493.1250000	34289241.5966871
## X53	-115111.2878664	0.0032202
## X54	-2209625971132224503808.0000000	2259369896535990.0000000
## X55	-37517750018965.8671875	9721379.4010338
## X56	605199811146251239424.0000000	618823786425694.3750000
## X57	-1377838456967522.5000000	10559349.2127029
## X58	-416671.9369221	0.0067287
## X59	-133962.9149332	0.0018808
## X60	-1970983876970873356288.0000000	2015352757388043.0000000
## X61	3456621710105407.0000000	86105188.2379008
## X62	-20467877072612.1250000	84167887.4097392
## X63	-15126.8207547	0.0015570
## X64	164031833095332.2187500	13773798.5227675
## X65	-5596858581754939.0000000	430285063.0177115
## X66	2869397278865996.0000000	515611972.4494765
## X67	-132198.7820026	0.0015964
## X68	1625186356087748.5000000	67526247.5542146
## X69	-2149618684838555.0000000	81775877.8324365
## X70	-3523319003549844.5000000	58273454.8268475
## X71	8341.2562027	0.0004010
## X72	1430.8472679	0.0003307

## X73	-74108869683436304.0000000	1718597696.3550916
## X74	-21443.3707844	0.0003355
## X75	-8808095994031961.0000000	709063857.2721953
## X76	-153246393465485.2812500	6979420.6307562
## X82	-868977826373918.1250000	44627211.4787531
## X85	1383878271926223.7500000	35772501.6389378
## X86	-1201412510424466.0000000	59351605.5431496
## X87	-57991.6915018	0.0024935
## X88	1964706613133726.7500000	86332886.1440868
## X89	5958137359182219.0000000	89686543.6969836
## X90	-2298060225279485.5000000	125109021.3483346
## X91	-14294886339432660.0000000	492426317.5584345
## X93	1652645407011857.0000000	76495068.7825037
##	z value	Pr(> z)
## (Intercept)	1108834	<0.0000000000000002 ***
## X6	-29226250	<0.0000000000000002 ***
## X7	29194483	<0.0000000000000002 ***
## X8	49703176	<0.0000000000000002 ***
## X9	-29224991	<0.0000000000000002 ***
## X10	45919510	<0.0000000000000002 ***
## X11	-8728171	<0.0000000000000002 ***
## X12	17518963	<0.0000000000000002 ***
## X13	-58655021	<0.0000000000000002 ***
## X14	-131167341	<0.0000000000000002 ***
## X15	-126220452	<0.0000000000000002 ***
## X18	-32616963	<0.0000000000000002 ***
## X19	-96855187	<0.0000000000000002 ***
## X20	10975760	<0.0000000000000002 ***
## X21	39100662	<0.0000000000000002 ***
## X24	-5890968	<0.0000000000000002 ***
## X25	-18404611	<0.0000000000000002 ***
## X26	10267199	<0.0000000000000002 ***
## X27	-9165009	<0.0000000000000002 ***
## X28	1989830	<0.0000000000000002 ***
## X29	-8339408	<0.0000000000000002 ***
## X30	47158948	<0.0000000000000002 ***
## X31	14981471	<0.0000000000000002 ***
## X32	19967390	<0.0000000000000002 ***
## X33	-50817859	<0.0000000000000002 ***
## X34	14228939	<0.0000000000000002 ***
## X35	43567162	<0.0000000000000002 ***
## X36	35310670	<0.0000000000000002 ***
## X37	38552483	<0.0000000000000002 ***
## X39	-26154902	<0.0000000000000002 ***
## X40	32548149	<0.0000000000000002 ***
## X44	-9047662	<0.0000000000000002 ***
## X45	-39286875	<0.0000000000000002 ***
## X46	-232636868	<0.0000000000000002 ***
## X47	-34920973	<0.0000000000000002 ***
## X48	-7418175	<0.0000000000000002 ***
## X49	15755695	<0.0000000000000002 ***
## X50	-811534	<0.0000000000000002 ***
## X51	52049365	<0.0000000000000002 ***
## X52	18865840	<0.0000000000000002 ***

```
## X53      -35746949 <0.0000000000000002 ***
## X54      -977983  <0.0000000000000002 ***
## X55      -3859303 <0.0000000000000002 ***
## X56       977984  <0.0000000000000002 ***
## X57     -130485168 <0.0000000000000002 ***
## X58      -61924315 <0.0000000000000002 ***
## X59      -71227328 <0.0000000000000002 ***
## X60      -977985  <0.0000000000000002 ***
## X61      40144175  <0.0000000000000002 ***
## X62      -243179  <0.0000000000000002 ***
## X63      -9715390 <0.0000000000000002 ***
## X64      11908976 <0.0000000000000002 ***
## X65     -13007327 <0.0000000000000002 ***
## X66       5565032 <0.0000000000000002 ***
## X67     -82812341 <0.0000000000000002 ***
## X68       24067476 <0.0000000000000002 ***
## X69     -26286709 <0.0000000000000002 ***
## X70     -60461818 <0.0000000000000002 ***
## X71       20798640 <0.0000000000000002 ***
## X72       4326372 <0.0000000000000002 ***
## X73     -43121709 <0.0000000000000002 ***
## X74     -63907466 <0.0000000000000002 ***
## X75     -12422148 <0.0000000000000002 ***
## X76     -21956893 <0.0000000000000002 ***
## X82     -19471927 <0.0000000000000002 ***
## X85       38685532 <0.0000000000000002 ***
## X86     -20242292 <0.0000000000000002 ***
## X87     -23257188 <0.0000000000000002 ***
## X88       22757337 <0.0000000000000002 ***
## X89       66432902 <0.0000000000000002 ***
## X90     -18368461 <0.0000000000000002 ***
## X91     -29029493 <0.0000000000000002 ***
## X93       21604601 <0.0000000000000002 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 1501.9  on 5268  degrees of freedom
## Residual deviance: 11894.4  on 5196  degrees of freedom
## AIC: 12040
##
## Number of Fisher Scoring iterations: 25
```

*# we can see that all our variables are highly significant, this is not that
surprising since most of our data is financial/account related. Lets look at
some correlations*

variable correlations with deposit for our original data
cor(train_og)[1,]

```
##      bankrupt      X6      X7      X8      X9
## 1.000000000000 -0.00007453889 -0.00859526287 -0.00905603766 -0.01953743085
##           X10      X11      X12      X13      X14
```

```
## -0.00808300696 -0.00894448007 -0.02103260822 -0.07601010012 -0.02214680734
##          X15          X18          X19          X20          X21
## -0.11590293570 -0.17066675112 -0.22970143937 -0.06865124674 -0.00410877305
##          X24          X25          X26          X27          X28
##  0.00040277150 -0.01542160582 -0.03961266194 -0.03959341737 -0.01032684566
##          X29          X30          X31          X32          X33
## -0.04631284864  0.07432027893 -0.01644192193 -0.08173993269 -0.00251570132
##          X34          X35          X36          X37          X39
##  0.02851076149 -0.00830286100  0.01677766684  0.24621563363 -0.00831038112
##          X40          X44          X45          X46          X47
##  0.17969279455  0.07261501572 -0.06791619403 -0.00376523494 -0.00704420745
##          X48          X49          X50          X51          X52
## -0.01165696982  0.06444932724  0.01707013802  0.07545624030 -0.10047018078
##          X53          X54          X55          X56          X57
## -0.00686167084 -0.19806952513 -0.07828541138 -0.05397970860 -0.09246088276
##          X58          X59          X60          X61          X62
## -0.00355547355  0.04576263927  0.18871935813 -0.07669897197 -0.00287457809
##          X63          X64          X65          X66          X67
## -0.00175106921 -0.02398596102 -0.15599397455  0.16177220370 -0.01747418867
##          X68          X69          X70          X71          X72
## -0.22144477813 -0.00698721071  0.15430324129  0.01553398983  0.01477656233
##          X73          X74          X75          X76          X82
## -0.00313044000 -0.02269523483  0.00096476077  0.07128956421 -0.11485546262
##          X85          X86          X87          X88          X89
##  0.14077739128 -0.32748227827  0.02048078633 -0.00184035707 -0.10895366954
##          X90          X91          X93
## -0.19318297847  0.17344579137 -0.00873311758
```


```
# looking at the the variables correlated at 10%
cor(train_og)[1,abs(cor(train_og)[1,])>0.1]
```

```
## bankrupt      X15      X18      X19      X37      X40      X52
## 1.0000000 -0.1159029 -0.1706668 -0.2297014  0.2462156  0.1796928 -0.1004702
##          X54      X60      X65      X66      X68      X70      X82
## -0.1980695  0.1887194 -0.1559940  0.1617722 -0.2214448  0.1543032 -0.1148555
##          X85      X86      X89      X90      X91
##  0.1407774 -0.3274823 -0.1089537 -0.1931830  0.1734458
```

```
# heatmap
cormat <- round(cor(train_og),2)[,abs(cor(train_og)[1,])>0.1]
melted_cormat <- melt(cormat)
head(melted_cormat)
```

```
##      Var1      Var2 value
## 1 bankrupt bankrupt  1.00
## 2      X6 bankrupt  0.00
## 3      X7 bankrupt -0.01
## 4      X8 bankrupt -0.01
## 5      X9 bankrupt -0.02
## 6     X10 bankrupt -0.01
```

```
ggplot(data = melted_cormat, aes(x=Var1, y=Var2, fill=value)) +  
  geom_tile()
```



```
# variable correlations with deposit for our bootstrap data  
cor(train_boot)[1,]
```

##	bankrupt	X6	X7	X8	X9	X10
##	1.000000000	0.001269780	-0.014748590	-0.018572125	-0.032999701	-0.014278401
##	X11	X12	X13	X14	X15	X18
##	-0.030802183	-0.025822878	-0.168364959	-0.048692056	-0.225424492	-0.320522574
##	X19	X20	X21	X24	X25	X26
##	-0.437360527	-0.139624128	-0.006901642	0.002476036	-0.020831916	-0.043735480
##	X27	X28	X29	X30	X31	X32
##	-0.043137051	-0.012868486	-0.096389280	0.082855226	-0.030870279	-0.167379526
##	X33	X34	X35	X36	X37	X39
##	-0.006901642	0.022963415	-0.032722234	0.048076638	0.445076438	-0.011738453
##	X40	X44	X45	X46	X47	X48
##	0.211220586	0.133906319	-0.126361692	-0.009622971	-0.011547053	-0.019794676
##	X49	X50	X51	X52	X53	X54
##	0.130575635	0.027971782	0.058570741	-0.204513416	-0.017522941	-0.370352196
##	X55	X56	X57	X58	X59	X60
##	-0.157169733	-0.097343532	-0.189795499	-0.169431354	0.080571901	0.369217121
##	X61	X62	X63	X64	X65	X66
##	-0.159277479	0.012108297	0.007678266	-0.023290180	-0.167138928	0.186662988
##	X67	X68	X69	X70	X71	X72


```
## -0.033516710 -0.422099219 -0.260502138 0.283345302 0.020842594 0.009377460
##          X73          X74          X75          X76          X82          X85
## -0.014860074 -0.057185299 -0.008900479 0.111714515 -0.247827095 0.143673943
##          X86          X87          X88          X89          X90          X91
## -0.501051859 0.037906053 -0.008434189 -0.196022271 -0.196609362 0.191929935
##          X93
## -0.040979441
```


```
# looking at the the variables correlated at 10%
cor(train_boot)[1,abs(cor(train_boot)[1,])>0.1]
```

```
## bankrupt      X13      X15      X18      X19      X20      X32
## 1.0000000 -0.1683650 -0.2254245 -0.3205226 -0.4373605 -0.1396241 -0.1673795
##          X37      X40      X44      X45      X49      X52      X54
## 0.4450764 0.2112206 0.1339063 -0.1263617 0.1305756 -0.2045134 -0.3703522
##          X55      X57      X58      X60      X61      X65      X66
## -0.1571697 -0.1897955 -0.1694314 0.3692171 -0.1592775 -0.1671389 0.1866630
##          X68      X69      X70      X76      X82      X85      X86
## -0.4220992 -0.2605021 0.2833453 0.1117145 -0.2478271 0.1436739 -0.5010519
##          X89      X90      X91
## -0.1960223 -0.1966094 0.1919299
```

```
# heatmap
cormat1 <- round(cor(train_boot),2)[,abs(cor(train_boot)[1,])>0.1]
melted_cormat1 <- melt(cormat1)
head(melted_cormat1)
```

```
##      Var1      Var2 value
## 1 bankrupt bankrupt 1.00
## 2      X6 bankrupt 0.00
## 3      X7 bankrupt -0.01
## 4      X8 bankrupt -0.02
## 5      X9 bankrupt -0.03
## 6     X10 bankrupt -0.01
```

```
ggplot(data = melted_cormat1, aes(x=Var1, y=Var2, fill=value)) +
  geom_tile()
```



```
# variable correlations with deposit for our adas data
cor(train_adas)[1,]
```

```
##      bankrupt      X6      X7      X8      X9
##  1.0000000000  0.0001337597 -0.0266761346 -0.0260248922 -0.0591843779
##           X10      X11      X12      X13      X14
## -0.0244059690 -0.0277445105 -0.0892203619 -0.2565259865 -0.0858258286
##           X15      X18      X19      X20      X21
## -0.3722347901 -0.5240936782 -0.6059632367 -0.2371465250 -0.0162144854
##           X24      X25      X26      X27      X28
##  0.0004678217 -0.0606238577 -0.1168736839 -0.1171019097 -0.0387250396
##           X29      X30      X31      X32      X33
## -0.1529343306  0.0443583900 -0.0652131961 -0.1657137965 -0.0099284967
##           X34      X35      X36      X37      X39
##  0.0310383860 -0.0345936967  0.0326812695  0.5871421185 -0.0253004691
##           X40      X44      X45      X46      X47
##  0.1867017742  0.1105480089 -0.2006302899 -0.0218274455 -0.0277919855
##           X48      X49      X50      X51      X52
## -0.0673366054  0.1327373234  0.0661134419  0.0504894102 -0.3397749652
##           X53      X54      X55      X56      X57
## -0.0270723120 -0.5066096060 -0.2152955532 -0.1456141015 -0.3130036771
##           X58      X59      X60      X61      X62
## -0.0140314286  0.0559302513  0.4929489981 -0.2455596375 -0.0041206130
##           X63      X64      X65      X66      X67
## -0.0292786635 -0.0557531127 -0.1486698405  0.1588117980 -0.0688140422
##           X68      X69      X70      X71      X72
```

```
## -0.5243143065 -0.0275030368 0.3105872814 0.0312979905 0.0258742879
##          X73          X74          X75          X76          X82
## -0.0096233022 -0.0826320429 0.0037244488 0.1712961917 -0.3085143185
##          X85          X86          X87          X88          X89
## 0.1494133189 -0.5700673798 0.0306836340 -0.0057688397 -0.3350124767
##          X90          X91          X93
## -0.1719110436 0.1693506083 -0.0195169333
```


```
# looking at the the variables correlated at 10%
cor(train_adas)[1,abs(cor(train_adas)[1,])>0.1]
```

```
## bankrupt      X13      X15      X18      X19      X20      X26
## 1.0000000 -0.2565260 -0.3722348 -0.5240937 -0.6059632 -0.2371465 -0.1168737
##          X27      X29      X32      X37      X40      X44      X45
## -0.1171019 -0.1529343 -0.1657138 0.5871421 0.1867018 0.1105480 -0.2006303
##          X49      X52      X54      X55      X56      X57      X60
## 0.1327373 -0.3397750 -0.5066096 -0.2152956 -0.1456141 -0.3130037 0.4929490
##          X61      X65      X66      X68      X70      X76      X82
## -0.2455596 -0.1486698 0.1588118 -0.5243143 0.3105873 0.1712962 -0.3085143
##          X85      X86      X89      X90      X91
## 0.1494133 -0.5700674 -0.3350125 -0.1719110 0.1693506
```

```
# heatmap
cormat2 <- round(cor(train_adas),2)[,abs(cor(train_adas)[1,])>0.1]
melted_cormat2 <- melt(cormat2)
head(melted_cormat2)
```

```
##      Var1      Var2 value
## 1 bankrupt bankrupt 1.00
## 2      X6 bankrupt 0.00
## 3      X7 bankrupt -0.03
## 4      X8 bankrupt -0.03
## 5      X9 bankrupt -0.06
## 6     X10 bankrupt -0.02
```

```
ggplot(data = melted_cormat2, aes(x=Var1, y=Var2, fill=value)) +
  geom_tile()
```



*# we are mainly concerned with the variables which are correlated to bankrupt
 # at 0.1 and more, thus we only plot heatmaps for those variables, to make sure
 # our model doesnt not have any multi-collinearity. We can see that there is some
 # correlation, bit its not significant enough for us to worry about for now.
 # thus based on the above correlations we will use the common variables that are
 # correlated to bankrupt for all the 3 datasets, so we can have comparable
 # results. We have 18 variables that are common in the 3 training datasets.
 # 'X15 ', 'X18 ', 'X19 ', 'X37 ', 'X40 ', 'X52 ', 'X54 ', 'X60 ', 'X65 ', 'X66 ', 'X68 ', 'X70 ', 'X82 ', 'X85 ', 'X86 ', 'X89 ', 'X9*

```
train_og <- train_og[,c(1,11,12,13,29,31,40,42,48,53,54,56,58,65,66,67,70,71,72)]
train_boot <- train_boot[,c(1,11,12,13,29,31,40,42,48,53,54,56,58,65,66,67,70,71,72)]
train_adas <- train_adas[,c(1,11,12,13,29,31,40,42,48,53,54,56,58,65,66,67,70,71,72)]
```

since we have removed these columns for our training data we will match out test set as well.

```
dat.test <- dat.test[,c(1,11,12,13,29,31,40,42,48,53,54,56,58,65,66,67,70,71,72)]
```

Part 2
LOGISTIC REGRESSION

part 2a
on our original train data

```
glm_base <- glm(bankrupt ~ ., data=train_og, family='binomial')
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(glm_base)
```

```
##
## Call:
## glm(formula = bankrupt ~ ., family = "binomial", data = train_og)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
##    -8.49      0.00      0.00      0.00      8.49
##
## Coefficients:
##              Estimate      Std. Error  z value
## (Intercept) 22293475726326296    156320854   142613574
## X15         -3294058247162912      6937270  -474834949
## X18         -5137897431815598      47490858 -108187083
## X19         -5955934136660202      59829356 -99548691
## X37          2934067722166466      40654152   72171415
## X40          4898560537117784      248513369   19711457
## X52          1232179730472399      30628200   40230236
## X54          2313306120747400      22972669   100698186
## X60          1567632463229763      42684610   36725941
## X65         -21480204190973108      142568659 -150665681
## X66         -23834031343701252      315985035  -75427722
## X68           3230897396364499      64665002   49963617
## X70         -5133109221427403      48852581 -105073450
## X82          -523395175375647      18461454  -28350701
## X85           1148891185114548      29888395   38439374
## X86         -5545165800535981      54217060 -102277138
## X89           1681189375362367      71993150   23352074
## X90         -3136378855956487      118763290  -26408656
## X91           6409121791671571      402603737   15919181
##              Pr(>|z|)
## (Intercept) <0.0000000000000002 ***
## X15         <0.0000000000000002 ***
## X18         <0.0000000000000002 ***
## X19         <0.0000000000000002 ***
## X37         <0.0000000000000002 ***
## X40         <0.0000000000000002 ***
## X52         <0.0000000000000002 ***
## X54         <0.0000000000000002 ***
## X60         <0.0000000000000002 ***
## X65         <0.0000000000000002 ***
## X66         <0.0000000000000002 ***
## X68         <0.0000000000000002 ***
## X70         <0.0000000000000002 ***
## X82         <0.0000000000000002 ***
## X85         <0.0000000000000002 ***
## X86         <0.0000000000000002 ***
## X89         <0.0000000000000002 ***
## X90         <0.0000000000000002 ***
## X91         <0.0000000000000002 ***
```

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1501.9  on 5268  degrees of freedom
## Residual deviance: 13192.0  on 5250  degrees of freedom
## AIC: 13230
##
## Number of Fisher Scoring iterations: 25

# prediction using our train set for error
yhat.train.lr <- predict(glm_base, train_og,
                        type = "response")

# classifying the deposit variable as 0/1 or no/yes with a threshold of 0.5
lr_yhat.train.class <- ifelse(yhat.train.lr > 0.5, 1, 0)

# confusion matrix of our train set
tab.lr_train <- table(train_og$bankrupt,
                     lr_yhat.train.class,
                     dnn = c("Actual","Predicted"))
tab.lr_train

##      Predicted
## Actual0    1
##      0 5020   79
##      1  104   66

# overall train error
lr_train <- mean(lr_yhat.train.class != train_og$bankrupt)
lr_train

## [1] 0.03473145

# class 1 test error
lr_train_class1 <- tab.lr_train[2,1]/5269
lr_train_class1

## [1] 0.01973809

# class 0 test error
lr_train_class0 <- tab.lr_train[1,2]/5269
lr_train_class0

## [1] 0.01499336

# prediction using our test set for error
yhat.test.lr <- predict(glm_base, dat.test,
                      type = "response")
lr_yhat.test.class <- ifelse(yhat.test.lr > 0.5, 1, 0)
```

```
tab.lr_test <- table(dat.test$bankrupt,
                     lr_yhat.test.class,
                     dnn = c("Actual","Predicted"))
tab.lr_test
```

```
##          Predicted
## Actual    0      1
##          0 1476   24
##          1   36   14
```

```
# overall test error
lr_test <- mean(lr_yhat.test.class != dat.test$bankrupt)
lr_test
```

```
## [1] 0.03870968
```

```
# class 1 test error
lr_test_class1 <- tab.lr_test[2,1]/1550
lr_test_class1
```

```
## [1] 0.02322581
```

```
# class 0 test error
lr_test_class0 <- tab.lr_test[1,2]/1550
lr_test_class0
```

```
## [1] 0.01548387
```

```
# part 2c
# model on our train_boot data
glm_boot <- glm(bankrupt ~ ., data=train_boot, family='binomial')
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(glm_boot)
```

```
##
## Call:
## glm(formula = bankrupt ~ ., family = "binomial", data = train_boot)
##
## Deviance Residuals:
## Min       1Q   Median       3Q      Max
## -4.5761  -0.3503  -0.1807  -0.0726   3.0443
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  110.9584   32.2186   3.444  0.000573 ***
## X15          -0.2585    0.5855  -0.442  0.658845
## X18          -7.5216    5.5964  -1.344  0.178948
```

```
## X19      -52.7236      7.1637   -7.360 0.0000000000000184 ***
## X37      18.7193      3.7202    5.032 0.000000486048885 ***
## X40      88.4478     20.1373    4.392 0.000011218647184 ***
## X52       1.7424      3.2816    0.531      0.595450
## X54       2.3515      2.1816    1.078      0.281098
## X60       7.9126      4.4397    1.782      0.074709 .
## X65     -73.9842     15.0535   -4.915 0.000000888918430 ***
## X66    -107.0638     53.3659   -2.006      0.044833 *
## X68      -2.6829      4.0238   -0.667      0.504927
## X70      -4.2330      2.4200   -1.749      0.080256 .
## X82      -1.9469      1.3200   -1.475      0.140228
## X85      17.3475     387.1364    0.045      0.964259
## X86      -2.1190      4.1769   -0.507      0.611937
## X89      17.2668      5.8173    2.968      0.002996 **
## X90     -40.1819     21.1498   -1.900      0.057449 .
## X91     -74.9727     37.5379   -1.997      0.045798 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 2870.8  on 3499  degrees of freedom
## Residual deviance: 1542.8  on 3481  degrees of freedom
## AIC: 1580.8
##
## Number of Fisher Scoring iterations: 15
```

```
# prediction using our train set for error
yhat.train.lr_boot <- predict(glm_boot, train_boot,
                             type = "response")

# classifying the deposit variable as 0/1 or no/yes with a threshold of 0.5
lr_yhat.train.boot.class <- ifelse(yhat.train.lr_boot > 0.5, 1, 0)

# confusion matrix of our train set
tab.lr_train_boot <- table(train_boot$bankrupt,
                          lr_yhat.train.boot.class,
                          dnn = c("Actual","Predicted"))
tab.lr_train_boot
```

```
##      Predicted
## Actual0    1
##      0 2910   90
##      1  229  271
```

```
# overall train error
lr_train_boot <- mean(lr_yhat.train.boot.class != train_boot$bankrupt)
lr_train_boot
```

```
## [1] 0.09114286
```



```
# class 1 test error
lr_train_boot_class1 <- tab.lr_train_boot[2,1]/3500
lr_train_boot_class1
```

```
## [1] 0.06542857
```

```
# class 0 test error
lr_train_boot_class0 <- tab.lr_train_boot[1,2]/3500
lr_train_boot_class0
```

```
## [1] 0.02571429
```

```
# prediction using our test set for error
yhat.test.lr_boot <- predict(glm_boot, dat.test,
                             type = "response")
lr_yhat.test.boot.class <- ifelse(yhat.test.lr_boot > 0.5, 1, 0)

tab.lr_test.boot <- table(dat.test$bankrupt,
                          lr_yhat.test.boot.class,
                          dnn = c("Actual", "Predicted"))

tab.lr_test.boot
```

```
##      Predicted
## Actual    0    1
##      0 1457   43
##      1   30   20
```

```
# overall test error
lr_test_boot <- mean(lr_yhat.test.boot.class != dat.test$bankrupt)
lr_test_boot
```

```
## [1] 0.04709677
```

```
# class 1 test error
lr_test_boot_class1 <- tab.lr_test.boot[2,1]/1550
lr_test_boot_class1
```

```
## [1] 0.01935484
```

```
# class 0 test error
lr_test_boot_class0 <- tab.lr_test.boot[1,2]/1550
lr_test_boot_class0
```

```
## [1] 0.02774194
```

```
# part 2d
# on our train_adas data
glm_adas <- glm(bankrupt ~ ., data=train_adas, family='binomial')
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(glm_adas)
```

```
##
## Call:
## glm(formula = bankrupt ~ ., family = "binomial", data = train_adas)
##
## Deviance Residuals:
## Min       1Q   Median       3Q      Max
## -8.4904  -0.3372   0.0000   0.3738   2.5836
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  167.0427    19.5686   8.536 < 0.0000000000000002 ***
## X15           0.5194     0.2743   1.893    0.05829 .
## X18          -4.1505     2.9049  -1.429    0.15306
## X19         -61.3436     4.4711 -13.720 < 0.0000000000000002 ***
## X37           16.9427     2.2456   7.545  0.00000000000000453 ***
## X40          184.8304    12.5566  14.720 < 0.0000000000000002 ***
## X52          -2.2390     2.0327  -1.101    0.27069
## X54           9.7153     1.5556   6.245  0.0000000004231045 ***
## X60           16.5366     2.7843   5.939  0.0000000028618956 ***
## X65          -144.4739    15.0711  -9.586 < 0.0000000000000002 ***
## X66          -189.8534    38.7263  -4.902  0.0000009465525807 ***
## X68           -0.3500     2.8739  -0.122    0.90306
## X70           -8.0730     1.3768  -5.864  0.0000000045275152 ***
## X82           -1.8411     0.7157  -2.572    0.01010 *
## X85           10.6461     2.6403   4.032  0.0000552727583670 ***
## X86          -14.1992     2.9749  -4.773  0.0000018146094624 ***
## X89           6.7979     2.8602   2.377    0.01747 *
## X90          -35.2725    11.9687  -2.947    0.00321 **
## X91          -95.7023    34.5691  -2.768    0.00563 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 14210.6  on 10250  degrees of freedom
## Residual deviance:  5744.3  on 10232  degrees of freedom
## AIC: 5782.3
##
## Number of Fisher Scoring iterations: 10
```

```
# prediction using our train set for error
```

```
yhat.train.lr_adas <- predict(glm_adas, train_adas,
                              type = "response")
```

```
# classifying the deposit variable as 0/1 or no/yes with a threshold of 0.5
```

```
lr_yhat.train.adas.class <- ifelse(yhat.train.lr_adas > 0.5, 1, 0)
```

```
# confusion matrix of our train set
```

```
tab.lr_train_adas <- table(train_adas$bankrupt,
                           lr_yhat.train.adas.class,
```

```

                                dnn = c("Actual","Predicted"))
tab.lr_train_adas

##          Predicted
## Actual    0     1
##          0 4524  575
##          1  533 4619

# overall train error
lr_train_adas <- mean(lr_yhat.train.adas.class != train_adas$bankrupt)
lr_train_adas

```

```
## [1] 0.108087
```

```

# class 1 test error
lr_train_adas_class1 <- tab.lr_train_adas[2,1]/10251
lr_train_adas_class1

```

```
## [1] 0.05199493
```

```

# class 0 test error
lr_train_adas_class0 <- tab.lr_train_adas[1,2]/10251
lr_train_adas_class0

```

```
## [1] 0.05609209
```

```

# prediction using our test set for error
yhat.test.lr_adas <- predict(glm_adas, dat.test,
                             type = "response")
lr_yhat.test.adas.class <- ifelse(yhat.test.lr_adas > 0.5, 1, 0)

tab.lr_test_adas <- table(dat.test$bankrupt,
                          lr_yhat.test.adas.class,
                          dnn = c("Actual","Predicted"))
tab.lr_test_adas

```

```

##          Predicted
## Actual    0     1
##          0 1335  165
##          1   10   40

```

```

# overall test error
lr_test_adas <- mean(lr_yhat.test.adas.class != dat.test$bankrupt)
lr_test_adas

```

```
## [1] 0.1129032
```

```
# class 1 test error
lr_test_adas_class1 <- tab.lr_test.adas[2,1]/1550
lr_test_adas_class1
```

```
## [1] 0.006451613
```

```
# class 0 test error
lr_test_adas_class0 <- tab.lr_test.adas[1,2]/1550
lr_test_adas_class0
```

```
## [1] 0.1064516
```

```
# After performing the predictions on the test dataset using the models built using the 3 train sets,
# the logistic regression performs the best[considering only performance]
```

```
# Part 3
# DECISION TREES
```

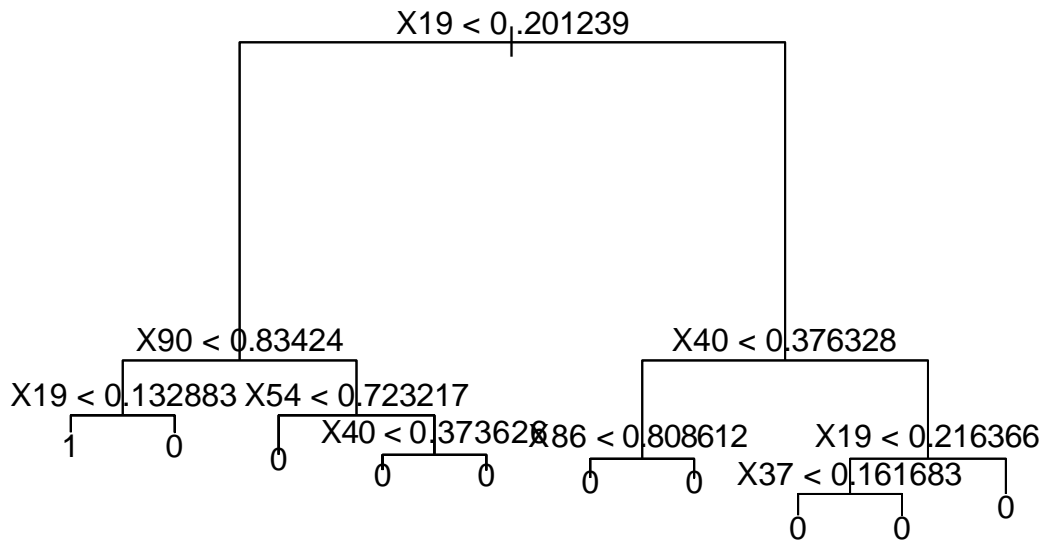
```
# part 3a
# on our original train data
# converting the Y-variable to a "factor" since classification tree models assume
# that the Y variable is qualitative
new.train_og <- train_og
new.train_og[,1] <- as.factor(new.train_og[,1])
```

```
# preparing our test data
new.dat.test <- dat.test
new.dat.test[,1] <- as.factor(new.dat.test[,1])
```

```
# building our first tree
tree_og <- tree(bankrupt ~., data = new.train_og)
summary(tree_og)
```

```
##
## Classification tree:
## tree(formula = bankrupt ~ ., data = new.train_og)
## Variables actually used in tree construction:
## [1] "X19" "X90" "X54" "X40" "X86" "X37"
## Number of terminal nodes: 10
## Residual mean deviance: 0.1634 = 859.4 / 5259
## Misclassification error rate: 0.03056 = 161 / 5269
```

```
# plotting our tree
plot(tree_og)
text(tree_og, pretty = 0)
```



```
# metadata
tree_og
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 5269 1502.00 0 ( 0.967736 0.032264 )
##    2) X19 < 0.201239 555 563.70 0 ( 0.794595 0.205405 )
##      4) X90 < 0.83424 186 249.20 0 ( 0.607527 0.392473 )
##        8) X19 < 0.132883 9 0.00 1 ( 0.000000 1.000000 ) *
##        9) X19 > 0.132883 177 231.60 0 ( 0.638418 0.361582 ) *
##      5) X90 > 0.83424 369 257.40 0 ( 0.888889 0.111111 )
##        10) X54 < 0.723217 49 66.92 0 ( 0.571429 0.428571 ) *
##        11) X54 > 0.723217 320 149.60 0 ( 0.937500 0.062500 )
##          22) X40 < 0.373626 138 11.85 0 ( 0.992754 0.007246 ) *
##          23) X40 > 0.373626 182 121.80 0 ( 0.895604 0.104396 ) *
##    3) X19 > 0.201239 4714 607.80 0 ( 0.988120 0.011880 )
##      6) X40 < 0.376328 3712 126.40 0 ( 0.997575 0.002425 )
##        12) X86 < 0.808612 1233 106.50 0 ( 0.992701 0.007299 ) *
##        13) X86 > 0.808612 2479 0.00 0 ( 1.000000 0.000000 ) *
##      7) X40 > 0.376328 1002 379.40 0 ( 0.953094 0.046906 )
##        14) X19 < 0.216366 306 213.50 0 ( 0.888889 0.111111 )
##          28) X37 < 0.161683 145 36.61 0 ( 0.972414 0.027586 ) *
##          29) X37 > 0.161683 161 154.80 0 ( 0.813665 0.186335 ) *
##        15) X19 > 0.216366 696 129.20 0 ( 0.981322 0.018678 ) *
```

```
# making predictions on our train data using tree1
tree.pred.train_og <- predict(tree_og, new.train_og, type = "class")
```

```
# confusion matrix
tab_tree_og <- table(new.train_og$bankrupt, tree.pred.train_og,
                     dnn = c("Actual", "Predicted"))
tab_tree_og
```

```
##      Predicted
## Actual    0    1
##      0 5099    0
##      1  161    9
```

```
# overall error
tree_og_err <- mean(new.train_og$bankrupt != tree.pred.train_og)
tree_og_err
```

```
## [1] 0.03055608
```

```
# class 1 train error
tree_og_class1 <- tab_tree_og[2,1]/5269
tree_og_class1
```

```
## [1] 0.03055608
```

```
# class 0 train error
tree_og_class0 <- tab_tree_og[1,2]/5269
tree_og_class0
```

```
## [1] 0
```

```
# making predictions using our test data
tree.pred.test <- predict(tree_og, new.dat.test, type = "class")
```

```
# confusion matrix
tab_tree_og.test <- table(new.dat.test$bankrupt, tree.pred.test,
                         dnn = c("Actual", "Predicted"))
tab_tree_og.test
```

```
##      Predicted
## Actual    0    1
##      0 1495    5
##      1   50    0
```

```
# overall error
tree_og_err.test <- mean(new.dat.test$bankrupt != tree.pred.test)
tree_og_err.test
```

```
## [1] 0.03548387
```

```
# class 1 test error
tree_og_class1.test <- tab_tree_og.test[2,1]/1550
tree_og_class1.test
```

```
## [1] 0.03225806
```


```
# class 0 test error
tree_og_class0.test <- tab_tree_og.test[1,2]/1550
tree_og_class0.test
```

```
## [1] 0.003225806
```


```
# pruning our decision tree
prune_og <- prune.misclass(tree_og)
names(prune_og)
```

```
## [1] "size" "dev" "k" "method"
```

```
# Plotting the results of the prune to identify the right tree size
plot(prune_og)
```




```
plot(prune_og$size, prune_og$dev, xlab = "Size of Tree",
     ylab = "Deviation/Deviance")
```



```
# pruning our tree1
prune.tree_og <- prune.misclass(tree_og, best = 4)
summary(prune.tree_og)
```

```
##
## Classification tree:
## snip.tree(tree = tree_og, nodes = c(3L, 5L))
## Variables actually used in tree construction:
## [1] "X19" "X90"
## Number of terminal nodes: 4
## Residual mean deviance: 0.2083 = 1097 / 5265
## Misclassification error rate: 0.03056 = 161 / 5269
```

```
# plotting our pruned tree
plot(prune.tree_og)
text(prune.tree_og, pretty = 0)
```

```
# making predictions on our test data using prune.tree1
pt1.pred <- predict(prune.tree_og, new.train_og, type = "class")
```

```
# confusion matrix
tab_pt1 <- table(new.train_og$bankrupt, pt1.pred,
                 dnn = c("Actual", "Predicted"))
```

```
tab_pt1
```

```
##      Predicted
## Actual    0    1
##      0 5099    0
##      1  161    9
```

```
# overall error
pt1_err <- mean(new.train_og$bankrupt != pt1.pred)
pt1_err
```

```
## [1] 0.03055608
```

```
# class 1 test error
pt1_class1 <- tab_pt1[2,1]/5269
pt1_class1
```

```
## [1] 0.03055608
```

```
# class 0 test error
pt1_class0 <- tab_pt1[1,2]/5269
pt1_class0
```

```
## [1] 0
```

```
# prediction using our test set for error
pt1.test.pred <- predict(prune.tree_og, new.dat.test, type = "class")
```

```
# confusion matrix
tab_pt1.test <- table(new.dat.test$bankrupt, pt1.test.pred,
                      dnn = c("Actual", "Predicted"))
```

```
tab_pt1.test
```

```
##      Predicted
## Actual    0    1
##      0 1495    5
##      1   50    0
```

```
# overall error
pt1_err.test <- mean(new.dat.test$bankrupt != pt1.test.pred)
pt1_err.test
```

```
## [1] 0.03548387
```

```
# class 1 test error
pt1_class1.test <- tab_pt1.test[2,1]/1550
pt1_class1.test
```

```
## [1] 0.03225806
```

```
# class 0 test error
pt1_class0.test <- tab_pt1.test[1,2]/1550
pt1_class0.test
```


```
## [1] 0.003225806
```

```
# part 3b
# model on our train_boot data
# converting the Y-variable to a "factor" since classification tree models assume
# that the Y variable is qualitative
new.train_boot <- train_boot
new.train_boot[,1] <- as.factor(new.train_boot[,1])

# building our first tree
tree_boot <- tree(bankrupt ~., data = new.train_boot)
summary(tree_boot)
```

```
##
## Classification tree:
## tree(formula = bankrupt ~ ., data = new.train_boot)
## Variables actually used in tree construction:
## [1] "X19" "X37" "X18" "X40" "X86" "X68" "X89" "X15"
## Number of terminal nodes: 13
## Residual mean deviance: 0.3419 = 1192 / 3487
## Misclassification error rate: 0.08057 = 282 / 3500
```

```
# plotting our tree
plot(tree_boot)
text(tree_boot, pretty = 0)
```



```
# metadata
tree_boot
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 3500 2871.00 0 ( 0.85714 0.14286 )
##    2) X19 < 0.210693 959 1312.00 0 ( 0.56726 0.43274 )
##      4) X37 < 0.161139 502 504.10 0 ( 0.79880 0.20120 )
##        8) X18 < 0.159854 155 214.40 0 ( 0.52903 0.47097 )
##          16) X19 < 0.199489 117 157.60 1 ( 0.40171 0.59829 ) *
##            17) X19 > 0.199489 38 20.99 0 ( 0.92105 0.07895 ) *
```

```
##          9) X18 > 0.159854 347  194.60 0 ( 0.91931 0.08069 )
##          18) X40 < 0.374025 192    0.00 0 ( 1.00000 0.00000 ) *
##          19) X40 > 0.374025 155  146.40 0 ( 0.81935 0.18065 ) *
##          5) X37 > 0.161139 457  568.00 1 ( 0.31291 0.68709 )
##          10) X86 < 0.728707 135  109.70 1 ( 0.14074 0.85926 ) *
##          11) X86 > 0.728707 322  429.20 1 ( 0.38509 0.61491 ) *
##          3) X19 > 0.210693 2541  744.70 0 ( 0.96655 0.03345 )
##          6) X68 < 0.938771 1121  601.90 0 ( 0.92417 0.07583 )
##          12) X40 < 0.382751 1006  367.60 0 ( 0.95527 0.04473 )
##          24) X40 < 0.372197 380    0.00 0 ( 1.00000 0.00000 ) *
##          25) X40 > 0.372197 626  323.60 0 ( 0.92812 0.07188 )
##          50) X89 < 0.616789 567  179.60 0 ( 0.96296 0.03704 ) *
##          51) X89 > 0.616789 59   79.73 0 ( 0.59322 0.40678 )
##          102) X86 < 0.80251 30   30.02 1 ( 0.20000 0.80000 )
##          204) X15 < 0.0121976 6    0.00 0 ( 1.00000 0.00000 ) *
##          205) X15 > 0.0121976 24    0.00 1 ( 0.00000 1.00000 ) *
##          103) X86 > 0.80251 29    0.00 0 ( 1.00000 0.00000 ) *
##          13) X40 > 0.382751 115  148.60 0 ( 0.65217 0.34783 ) *
##          7) X68 > 0.938771 1420    0.00 0 ( 1.00000 0.00000 ) *
```

```
# making predictions on our test data using tree1
tree.pred.train_boot <- predict(tree_boot, new.train_boot, type = "class")
```

```
# confusion matrix
tab_tree_boot <- table(new.train_boot$bankrupt, tree.pred.train_boot,
                      dnn = c("Actual", "Predicted"))
tab_tree_boot
```

```
##      Predicted
## Actual    0    1
##      0 2810  190
##      1   92  408
```

```
# overall error
tree_boot_err <- mean(new.train_boot$bankrupt != tree.pred.train_boot)
tree_boot_err
```

```
## [1] 0.08057143
```

```
# class 1 test error
tree_boot_class1 <- tab_tree_boot[2,1]/3500
tree_boot_class1
```

```
## [1] 0.02628571
```

```
# class 0 test error
tree_boot_class0 <- tab_tree_boot[1,2]/3500
tree_boot_class0
```

```
## [1] 0.05428571
```

```
# predictions on our test data
tree.pred.test_boot <- predict(tree_boot, new.dat.test, type = "class")
```

```
# confusion matrix
tab_tree_boot.test <- table(new.dat.test$bankrupt, tree.pred.test_boot,
                             dnn = c("Actual", "Predicted"))
tab_tree_boot.test
```

```
##      Predicted
## Actual    0    1
##      0 1419   81
##      1   19   31
```

```
# overall error
tree_boot_err.test <- mean(new.dat.test$bankrupt != tree.pred.test_boot)
tree_boot_err.test
```

```
## [1] 0.06451613
```

```
# class 1 test error
tree_boot_class1.test <- tab_tree_boot.test[2,1]/1550
tree_boot_class1.test
```

```
## [1] 0.01225806
```


```
# class 0 test error
tree_boot_class0.test <- tab_tree_boot.test[1,2]/1550
tree_boot_class0.test
```

```
## [1] 0.05225806
```


```
# pruning our decision tree
prune_boot <- prune.misclass(tree_boot)
names(prune_boot)
```

```
## [1] "size"    "dev"     "k"       "method"
```

```
# Plotting the results of the prune to identify the right tree size
plot(prune_boot)
```




```
plot(prune_boot$size, prune_boot$dev, xlab = "Size of Tree",  
     ylab = "Deviation/Deviance")
```



```
# pruning our tree1
prune.tree_boot <- prune.misclass(tree_boot, best = 11)
summary(prune.tree_boot)
```

```
##
## Classification tree:
## snip.tree(tree = tree_boot, nodes = c(5L, 9L))
## Variables actually used in tree construction:
## [1] "X19" "X37" "X18" "X68" "X40" "X89" "X86" "X15"
## Number of terminal nodes: 11
## Residual mean deviance: 0.3639 = 1269 / 3489
## Misclassification error rate: 0.08057 = 282 / 3500
```

```
# plotting our pruned tree
plot(prune.tree_boot)
text(prune.tree_boot, pretty = 0)
```



```
# making predictions on our test data using prune.tree1
pt2.pred <- predict(prune.tree_boot, new.train_boot, type = "class")
```

```
# confusion matrix
tab_pt2 <- table(new.train_boot$bankrupt, pt2.pred,
                 dnn = c("Actual", "Predicted"))
```

```
tab_pt2
```

```
##      Predicted
## Actual    0    1
##      0 2810  190
##      1   92  408
```

```
# overall error
pt2_err <- mean(new.train_boot$bankrupt != pt2.pred)
pt2_err
```

```
## [1] 0.08057143
```

```
# class 1 test error
pt2_class1 <- tab_pt2[2,1]/3500
pt2_class1
```

```
## [1] 0.02628571
```



```
# class 0 test error
pt2_class0 <- tab_pt2[1,2]/3500
pt2_class0
```

```
## [1] 0.05428571
```

```
# predictions using our test data
pt2.test.pred <- predict(prune.tree_boot, new.dat.test, type = "class")

# confusion matrix
tab_pt2.test <- table(new.dat.test$bankrupt, pt2.test.pred,
                      dnn = c("Actual", "Predicted"))

tab_pt2.test
```

```
##      Predicted
## Actual    0    1
##      0 1419   81
##      1   19   31
```

```
# overall error
pt2_err.test <- mean(new.dat.test$bankrupt != pt2.test.pred)
pt2_err.test
```

```
## [1] 0.06451613
```

```
# class 1 test error
pt2_class1.test <- tab_pt2.test[2,1]/1550
pt2_class1.test
```

```
## [1] 0.01225806
```

```
# class 0 test error
pt2_class0.test <- tab_pt2.test[1,2]/1550
pt2_class0.test
```

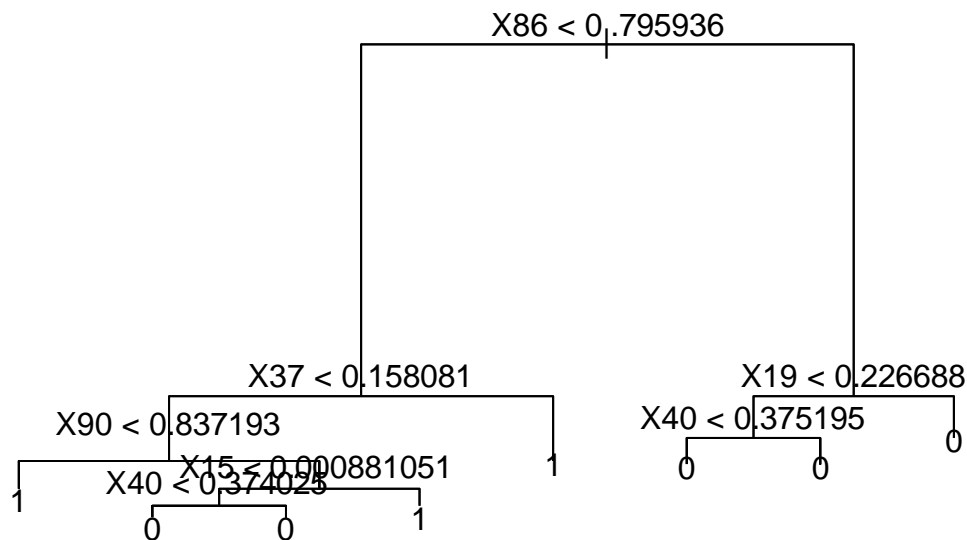
```
## [1] 0.05225806
```

```
# part 3c
# model on our train_adas data
# converting the Y-variable to a "factor" since classification tree models assume
# that the Y variable is qualitative
new.train_adas <- train_adas
new.train_adas[,1] <- as.factor(new.train_adas[,1])

# building our first tree
tree_adas <- tree(bankrupt ~., data = new.train_adas)
summary(tree_adas)
```

```
##
## Classification tree:
## tree(formula = bankrupt ~ ., data = new.train_adas)
## Variables actually used in tree construction:
## [1] "X86" "X37" "X90" "X15" "X40" "X19"
## Number of terminal nodes: 8
## Residual mean deviance: 0.5243 = 5370 / 10240
## Misclassification error rate: 0.1057 = 1084 / 10251
```

```
# plotting our tree
plot(tree_adas)
text(tree_adas, pretty = 0)
```



```
# metadata
tree_adas
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 10251 14210.0 1 ( 0.497415 0.502585 )
##    2) X86 < 0.795936 5759 5509.0 1 ( 0.184581 0.815419 )
##      4) X37 < 0.158081 1808 2479.0 1 ( 0.438606 0.561394 )
##        8) X90 < 0.837193 904 894.1 1 ( 0.195796 0.804204 ) *
##        9) X90 > 0.837193 904 1131.0 0 ( 0.681416 0.318584 )
##      18) X15 < 0.000881051 797 854.0 0 ( 0.772898 0.227102 )
```

```
##          36) X40 < 0.374025 419   148.5 0 ( 0.957041 0.042959 ) *
##          37) X40 > 0.374025 378   516.8 0 ( 0.568783 0.431217 ) *
##          19) X15 > 0.000881051 107     0.0 1 ( 0.000000 1.000000 ) *
##          5) X37 > 0.158081 3951  1970.0 1 ( 0.068337 0.931663 ) *
##          3) X86 > 0.795936 4492  2950.0 0 ( 0.898486 0.101514 )
##          6) X19 < 0.226688 2123  2196.0 0 ( 0.787565 0.212435 )
##          12) X40 < 0.375195 1180   491.8 0 ( 0.946610 0.053390 ) *
##          13) X40 > 0.375195 943   1278.0 0 ( 0.588547 0.411453 ) *
##          7) X19 > 0.226688 2369    71.6 0 ( 0.997889 0.002111 ) *
```

```
# making predictions on our test data using tree1
tree.pred.train_adas <- predict(tree_adas, new.train_adas, type = "class")
```

```
# confusion matrix
tab_tree_adas <- table(new.train_adas$bankrupt, tree.pred.train_adas,
                      dnn = c("Actual", "Predicted"))
tab_tree_adas
```

```
##          Predicted
## Actual      0      1
##          0 4652  447
##          1  637 4515
```

```
# overall error
tree_adas_err <- mean(new.train_adas$bankrupt != tree.pred.train_adas)
tree_adas_err
```

```
## [1] 0.1057458
```

```
# class 1 test error
tree_adas_class1 <- tab_tree_adas[2,1]/10251
tree_adas_class1
```

```
## [1] 0.06214028
```

```
# class 0 test error
tree_adas_class0 <- tab_tree_adas[1,2]/10251
tree_adas_class0
```

```
## [1] 0.0436055
```

```
# predictions using our test data
tree.pred.test_adas <- predict(tree_adas, new.dat.test, type = "class")
```

```
# confusion matrix
tab_tree_adas.test <- table(new.dat.test$bankrupt, tree.pred.test_adas,
                          dnn = c("Actual", "Predicted"))
tab_tree_adas.test
```

```
##          Predicted
## Actual      0      1
##          0 1381  119
##          1   17   33
```

```
# overall error
tree_adas_err.test <- mean(new.dat.test$bankrupt != tree.pred.test_adas)
tree_adas_err.test
```

```
## [1] 0.08774194
```

```
# class 1 test error
tree_adas_class1.test <- tab_tree_adas.test[2,1]/1550
tree_adas_class1.test
```

```
## [1] 0.01096774
```


```
# class 0 test error
tree_adas_class0.test <- tab_tree_adas.test[1,2]/1550
tree_adas_class0.test
```

```
## [1] 0.07677419
```


```
# pruning our decision tree
prune_adas <- prune.misclass(tree_adas)
names(prune_boot)
```

```
## [1] "size" "dev" "k" "method"
```

```
# Plotting the results of the prune to identify the right tree size
plot(prune_adas)
```




```
plot(prune_adas$size, prune_adas$dev, xlab = "Size of Tree",
     ylab = "Deviation/Deviance")
```



```
# pruning our tree1
prune.tree_adas <- prune.misclass(tree_adas, best = 5)
summary(prune.tree_adas)
```

```
##
## Classification tree:
## snip.tree(tree = tree_adas, nodes = c(3L, 18L))
## Variables actually used in tree construction:
## [1] "X86" "X37" "X90" "X15"
## Number of terminal nodes: 5
## Residual mean deviance: 0.6508 = 6668 / 10250
## Misclassification error rate: 0.1057 = 1084 / 10251
```

```
# plotting our pruned tree
plot(prune.tree_adas)
text(prune.tree_adas, pretty = 0)
```



```
# making predictions on our test data using prune.tree1
pt3.pred <- predict(prune.tree_adas, new.train_adas, type = "class")
```

```
# confusion matrix
tab_pt3 <- table(new.train_adas$bankrupt, pt3.pred,
                 dnn = c("Actual", "Predicted"))
```

```
tab_pt3
```

```
##      Predicted
## Actual    0    1
##      0 4652  447
##      1  637 4515
```

```
# overall error
pt3_err <- mean(new.train_adas$bankrupt != pt3.pred)
pt3_err
```

```
## [1] 0.1057458
```

```
# class 1 test error
pt3_class1 <- tab_pt3[2,1]/10251
pt3_class1
```

```
## [1] 0.06214028
```

```
# class 0 test error
pt3_class0 <- tab_pt3[1,2]/10251
pt3_class0
```

```
## [1] 0.0436055
```

```
# prediction using our test set for error
pt3.pred.test <- predict(prune.tree_adas, new.dat.test, type = "class")
```

```
# confusion matrix
tab_pt3.test <- table(new.dat.test$bankrupt, pt3.pred.test,
                      dnn = c("Actual", "Predicted"))
```

```
tab_pt3.test
```

```
##      Predicted
## Actual      0      1
##      0 1381   119
##      1   17    33
```

```
# overall error
pt3_err.test <- mean(new.dat.test$bankrupt != pt3.pred.test)
pt3_err.test
```

```
## [1] 0.08774194
```

```
# class 1 test error
pt3_class1.test <- tab_pt3.test[2,1]/1550
pt3_class1.test
```

```
## [1] 0.01096774
```

```
# class 0 test error
pt3_class0.test <- tab_pt3.test[1,2]/1550
pt3_class0.test
```

```
## [1] 0.07677419
```

```
# Part 4
# table to compare our train values
train.error <- data.frame(matrix(0,3,10))
names(train.error) <- c("Error", "Logistic Reg Train", "Logistic Reg Boot", "Logistic Reg ADAS",
                      "Decision Tree", "Pruned Tree", "Decision Tree Boot", "Pruned Tree Boot",
                      "Decision Tree Adas", "Pruned Tree ADAS")

train.error[1,] <- c('Overall Error', lr_train, lr_train_boot, lr_train_adas, tree_og_err, pt1_err,
                    tree_boot_err, pt2_err, tree_adas_err, pt3_err)
train.error[2,] <- c('Class 1 Err', lr_train_class1, lr_train_boot_class1, lr_train_adas_class1,
                    tree_og_class1, pt1_class1, tree_boot_class1, pt2_class1, tree_adas_class1,
                    pt3_class1)
```

```
train.error[3,] <- c('Class 0 Err', lr_train_class0, lr_train_boot_class0, lr_train_adas_class0,
                    tree_og_class0, pt1_class0, tree_boot_class0, pt2_class0, tree_adas_class0,
                    pt3_class0)
```

lets look at our table

```
train.error
```

```
##           Error Logistic Reg Train  Logistic Reg Boot  Logistic Reg ADAS
## 1 Overall Error 0.0347314480926172  0.0911428571428571  0.108087015900888
## 2   Class 1 Err 0.0197380907193016  0.0654285714285714  0.0519949273241635
## 3   Class 0 Err 0.0149933573733156  0.0257142857142857  0.0560920885767242
##           Decision Tree      Pruned Tree Decision Tree Boot  Pruned Tree Boot
## 1 0.0305560827481496 0.0305560827481496 0.0805714285714286 0.0805714285714286
## 2 0.0305560827481496 0.0305560827481496 0.0262857142857143 0.0262857142857143
## 3           0           0 0.0542857142857143 0.0542857142857143
##   Decision Tree Adas  Pruned Tree ADAS
## 1 0.105745780899424 0.105745780899424
## 2 0.062140278997171 0.062140278997171
## 3 0.0436055019022534 0.0436055019022534
```

table to compare our test values

```
test.error <- data.frame(matrix(0,3,10))
```

```
names(test.error) <- c("Error", "Logistic Reg", "Logistic Reg Boot", "Logistic Reg ADAS",
                      "Decision Tree", "Pruned Tree", "Decision Tree Boot", "Pruned Tree Boot",
                      "Decision Tree Adas", "Pruned Tree ADAS")
```

```
test.error[1,] <- c('Overall Error', lr_test, lr_test_boot, lr_test_adas, tree_og_err.test, pt1_err.test,
                    tree_boot_err.test, pt2_err.test, tree_adas_err.test, pt3_err.test)
```

```
test.error[2,] <- c('Class 1 Err', lr_test_class1, lr_test_boot_class1, lr_test_adas_class1,
                    tree_og_class1.test, pt1_class1.test, tree_boot_class1.test, pt2_class1.test,
                    , tree_adas_class1.test,
                    pt3_class1.test)
```

```
test.error[3,] <- c('Class 0 Err', lr_test_class0, lr_test_boot_class0, lr_test_adas_class0,
                    tree_og_class0.test, pt1_class0.test, tree_boot_class0.test, pt2_class0.test,
                    , tree_adas_class0.test,
                    pt3_class0.test)
```

```
test.error
```

```
##           Error      Logistic Reg  Logistic Reg Boot  Logistic Reg ADAS
## 1 Overall Error 0.0387096774193548  0.0470967741935484  0.112903225806452
## 2   Class 1 Err 0.0232258064516129  0.0193548387096774  0.00645161290322581
## 3   Class 0 Err 0.0154838709677419  0.027741935483871  0.106451612903226
##           Decision Tree      Pruned Tree Decision Tree Boot  Pruned Tree Boot
## 1 0.0354838709677419 0.0354838709677419 0.0645161290322581 0.0645161290322581
## 2 0.032258064516129 0.032258064516129 0.012258064516129 0.012258064516129
## 3 0.0032258064516129 0.0032258064516129 0.052258064516129 0.052258064516129
##   Decision Tree Adas  Pruned Tree ADAS
## 1 0.087741935483871 0.087741935483871
## 2 0.0109677419354839 0.0109677419354839
## 3 0.0767741935483871 0.0767741935483871
```



```
# Part 5
# Prediction of Bankruptcy for MediaTek INC. using our best model
```

```
max(dat.test$X40)
```

```
## [1] 0.4427657
```

```
# setting up our data
mediatek <- data.frame(X15 = 0.1262, X18 = 2.73, X19 = 0.7626, X37 = 0.2718,
                      X40 = 0.4427, X52 = 8.211, X54 = 0.2565, X60 = 0.2327,
                      X65 = 0.3523, X65 = 0.3523, X66 = 0.3195, X68 = 0.4985,
                      X70 = 0.7078, X82 = 0.2369, X85 = 0, X86 = 0.1942, X89 = 0.4666,
                      X90 = 0.2666, X91 = 0.3732)
```

```
# lets check
mediatek
```

```
##      X15  X18   X19   X37   X40  X52   X54   X60   X65  X65.1  X66
## 1 0.1262 2.73 0.7626 0.2718 0.4427 8.211 0.2565 0.2327 0.3523 0.3523 0.3195
##      X68   X70   X82 X85   X86   X89   X90   X91
## 1 0.4985 0.7078 0.2369  0 0.1942 0.4666 0.2666 0.3732
```

```
# prediction using our train set for error
yhat.mediatek <- predict(glm_base, mediatek,
                        type = "response")
```

```
# classifying the deposit variable as 0/1 or no/yes with a threshold of 0.5
lr_yhat.mediatek.class <- ifelse(yhat.mediatek > 0.5, 1, 0)
lr_yhat.mediatek.class
```

```
## 1
## 1
```

```
# we tried to predict the bankruptcy of Mediatek Inc. since their share prices
# have been dropping for the past year and eps is very low, However, our model did
# not do very well since it predicted that mediatek will go bankrupt, which is not
# likely since the company is doing financially well as of today. It must be mentioned
# that our model is not 100% accurate. Ths not included in the main presentation.
```